



Machine Learning


CS60050

Artificial Neural Networks



Neural networks


- ¬ Networks of processing units (neurons) with connections (synapses) between them
- ¬ Large number of neurons: 10^{10}
- ¬ Large connectivity: 10^5
- ¬ Parallel processing
- ¬ Distributed computation/memory
- ¬ Robust to noise, failures



Connectionism

- Alternative to *symbolism*
- Humans and evidence of connectionism/parallelism:
 - Physical structure of brain:
 - Neuron switching time: 10^{-3} second
- Complex, short-time computations:
 - Scene recognition time: 10^{-1} second
 - 100 inference steps doesn't seem like enough
 - much parallel computation
- Artificial Neural Networks (ANNs)
 - Many neuron-like threshold switching units
 - Many weighted interconnections among units
 - Highly parallel, distributed process
 - Emphasis on tuning weights automatically (search in weight space)

Biological neuron





Biological neuron

- **dendrites:** nerve fibres carrying electrical signals to the cell
- **cell body:** computes a non-linear function of its inputs
- **axon:** single long fiber that carries the electrical signal from the cell body to other neurons
- **synapse:** the point of contact between the axon of one cell and the dendrite of another, regulating a chemical connection whose strength affects the input to the cell.

Biological neuron

- A variety of different neurons exist (motor neuron, on-center off-surround visual cells...), with different branching structures
- The connections of the network and the strengths of the individual synapses establish the function of the network.

Biological inspiration



Biological inspiration

- The spikes travelling along the axon of the pre-synaptic neuron trigger the release of neurotransmitter substances at the synapse.
- The neurotransmitters cause excitation or inhibition in the dendrite of the post-synaptic neuron.
- The integration of the excitatory and inhibitory signals may produce spikes in the post-synaptic neuron.
- The contribution of the signals depends on the strength of the synaptic connection.

Hodgkin and Huxley model



- Hodgkin and Huxley experimented on squids and discovered how the signal is produced within the neuron
- This model was published in *Jour. of Physiology* (1952)
- They were awarded the 1963 Nobel Prize

When to consider ANNs

- Input is
 - ◆ high-dimensional
 - ◆ discrete or real-valued
 - ◆ e.g., raw sensor inputs
 - ◆ noisy
- *Long training times*
- Form of target function is unknown
- *Human readability is unimportant*
- Especially good for complex recognition problems
 - ◆ Speech recognition
 - ◆ Image classification
 - ◆ Financial prediction


Problems too hard to program

ALVINN: a perception system which learns to control the NAVLAB vehicles by watching a person drive



How many weights need to be learned?

Perceptron




- $-w_0$: threshold value or bias $\left(\sum_{i=1}^n w_i x_i\right) - (-w_0)$
- f (or $o()$) : activation function (thresholding unit), typically:
$$f(x) = \begin{cases} 1 & x > 0 \\ -1 & \text{otherwise} \end{cases}$$


Decision surface of a perceptron

- Decision surface is a hyperplane given by $\sum_{i=0}^n w_i x_i = 0$
- 2D case: the decision surface is a line
- Represents many useful functions: for example, $x_1 \wedge x_2$?

• $x_1 \wedge x_2$?





$x_1 \text{ XOR } x_2$?



Not linearly separable!

- Generalization to higher dimensions
 - ◆ Hyperplanes as decision surfaces

Learning Boolean AND



x_1	x_2	r
0	0	0
0	1	0
1	0	0
1	1	1

XOR

x_1	x_2	r
0	0	0
0	1	1
1	0	1
1	1	0

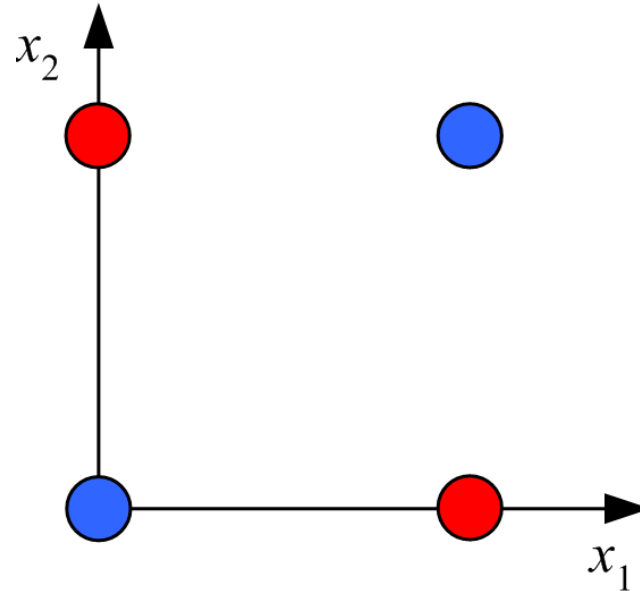
- No w_0, w_1, w_2 satisfy:

$$w_0 \leq 0$$

$$w_2 + w_0 > 0$$

$$w_1 + w_0 > 0$$

$$w_1 + w_2 + w_0 \leq 0$$



(Minsky and Papert, 1969)

Boolean functions

- Solution:
 - ◆ network of perceptrons
 - ◆ Any boolean function representable as DNF
 - ❖ 2 layers
 - ❖ Disjunction (layer 1) of conjunctions (layer 2)

- Example of XOR
 - ◆ $(X_1=1 \text{ AND } X_2=0) \text{ OR } (X_1=0 \text{ AND } X_2=1)$

- Practical problem of representing high-dimensional functions

Training rules

- Finding learning rules to build networks from TEs
- Will examine two major techniques
 - ◆ Perceptron training rule
 - ◆ Delta (gradient search) training rule (for more perceptrons as well as general ANNs)
- Both focused on learning weights
 - ◆ Hypothesis space can be viewed as set of weights

Perceptron training rule

- ITERATIVE RULE: $w_i := w_i + \Delta w_i$
 - ◆ where $\Delta w_i = \eta (t - o) x_i$
 - ◆ t is the target value
 - ◆ o is the perceptron output for x
 - ◆ η is small positive constant, called the **learning rate**
- Why rule works:
 - ◆ E.g., $t = 1, o = -1, x_i = 0.8, \eta = 0.1$
 - ◆ then $\Delta w_i = 0.16$ and $w_i x_i$ gets larger
 - ◆ o converges to t

Perceptron training rule

- The process will converge if
 - training data is linearly separable, and
 - η is sufficiently small
- But if the training data is not linearly separable, it may not converge (Minsky & Pappert)
 - Basis for Minsky/Pappert attack on NN approach
- Question: how to overcome problem:
 - different model of neuron?
 - different training rule?
 - both?

Gradient descent

- Solution: use alternate rule
 - ◆ More general
 - ◆ Basis for networks of units
 - ◆ Works in non-linearly separable cases
- Let $o(x) = w_0 + w_1x_1 + \dots + w_nx_n$
 - ◆ Simple example of linear unit (will generalize)
 - ◆ Omit the thresholding initially
- D is the set of training examples $\{d = \langle x, t_d \rangle\}$
- We will learn w_i 's that minimize the squared error

$$E[\vec{w}] \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$


Error minimization

- Look at error E as a function of weights $\{w_i\}$
- Slide down gradient of E in weight space
- Reach values of $\{w_i\}$ that correspond to minimum error
 - ◆ Look for global minimum
- Example of 2-dimensional case:
 - ◆ $E = w_1 * w_1 + w_2 * w_2$
 - ◆ Minimum at $w_1 = w_2 = 0$
- Look at general case of n -dimensional space of weights

Gradient descent

- Gradient “points” to the steepest increase:

$$\nabla E[\vec{w}] \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \dots, \frac{\partial E}{\partial w_n} \right]$$




- Training rule: $\Delta w = -\eta \nabla E[w]$
where η is a positive constant (learning rate)

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

Parabola with a single minima

- How might one interpret this update rule?

Gradient descent



$$\begin{aligned}\frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_{d \in D} \frac{\partial}{\partial w_i} (t_d - o_d)^2 \\ &= \frac{1}{2} \sum_{d \in D} 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d) \\ &= \sum_{d \in D} (t_d - o_d) \frac{\partial}{\partial w_i} (\vec{w}_d \cdot \vec{x}_d) \\ &= \sum_{d \in D} (t_d - o_d) (-x_{i,d})\end{aligned}$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i} = -\eta \sum_{d \in D} (t_d - o_d) (-x_{i,d}) = \eta \sum_{d \in D} (t_d - o_d) x_{i,d}$$

$$\Delta w_i = \sum_{d \in D} (\eta (t_d - o_d) x_{i,d})$$

Gradient descent algorithm

Gradient-Descent (training examples, η)

Each training example is a pair $\langle x, t \rangle$: x is the vector of input values, and t is the target output value. η is the learning rate (e.g., .05).

- Initialize each w_i to some small random value
- Repeat until the termination condition is met

1. *Initialize each Δw_i to zero*

2. *For each training example $\langle x, t \rangle$*

- ♦ *Input x to the unit and compute the output o*
- ♦ *For each linear unit weight w_i*

$$\Delta w_i \leftarrow \Delta w_i + \eta (t - o) x_i$$

3. *For each linear unit weight w_i*

$$w_i \leftarrow w_i + \Delta w_i$$

Also called

- LMS (Least Mean Square) rule
- Delta rule

- At each iteration, consider reducing η

Incremental (Stochastic) Gradient Descent

Batch mode Gradient Descent: $E_D[\vec{w}] \equiv \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$

- Repeat
 1. Compute the gradient $\nabla E_D[w]$
 2. $w \leftarrow w - \eta \nabla E_D[w]$

Incremental mode Gradient Descent: $E_d[\vec{w}] \equiv \frac{1}{2} (t_d - o_d)^2$

- Repeat
 - ◆ For each training example d in D
 1. Compute the gradient $\nabla E_d[w]$
 2. $w \leftarrow w - \eta \nabla E_d[w]$
- Incremental can approximate batch if η is small enough

Incremental Gradient Descent Algorithm

Incremental-Gradient-Descent (training examples, η)

Each training example is a pair $\langle x, t \rangle$: x is the vector of input values, and t is the target output value. η is the learning rate (e.g., .05).

- Initialize each w_i to some small random value
- Repeat until the termination condition is met
 1. *Initialize each Δw_i to zero*
 2. *For each $\langle x, t \rangle$*
 - ♦ *Input x to the unit and compute output o*
 - ♦ *For each linear unit weight w_i*
$$w_i \leftarrow w_i + \eta (t - o) x_i$$

Perceptron vs. Delta rule training

- Perceptron training rule guaranteed to succeed if
 - ◆ Training examples are linearly separable
 - ◆ Sufficiently small learning rate
- Delta training rule uses gradient descent
 - ◆ Guaranteed to converge to hypothesis with minimum squared error
 - ◆ Given sufficiently small learning rate
 - ◆ Even when training data contains noise
 - ◆ Even when training data not linearly separable
- Can generalize linear units to units with threshold
 - ◆ Just threshold the results


Perceptron vs. Delta rule training

- Delta/perceptron training rules appear same *but*
 - ◆ Perceptron rule trains discontinuous units
 - ◆ Guaranteed to converge under limited conditions
 - ◆ May not converge in general
 - ◆ Gradient rules trains over continuous response (unthresholded outputs)
 - ◆ Gradient rule always converges
 - Even with noisy training data
 - Even with non-separable training data
 - ◆ Gradient descent generalizes to other continuous responses
 - ◆ Can train perceptron with LMS rule
 - ◆ get prediction by thresholding outputs


Multilayer networks of sigmoid units

- Needed for relatively complex (i.e., typical) functions
- Want non-linear response units in many systems
 - ◆ Example (next slide) of phoneme recognition
 - ◆ Cascaded nets of linear units only give linear response
 - ◆ Sigmoid unit as example of many possibilities
- Want differentiable functions of weights
 - ◆ So can apply gradient descent
 - ◆ Minimization of error function
 - ◆ Step function perceptrons non-differentiable

Speech recognition example

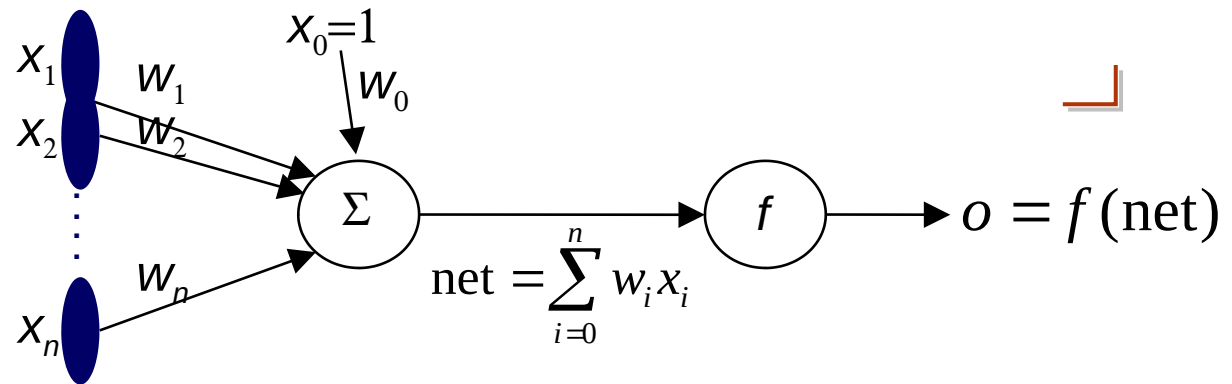


Multilayer networks



- Can have more than one hidden layer

Sigmoid unit



- f is the sigmoid function $f(x) = \frac{1}{1 + e^{-x}}$
- Derivative can be easily computed: $\frac{df(x)}{dx} = f(x)(1 - f(x))$
- Logistic equation
 - used in many applications
 - other functions possible (tanh)
- Single unit:
 - apply gradient descent rule
- Multilayer networks: **backpropagation**

Error Gradient for a Sigmoid Unit

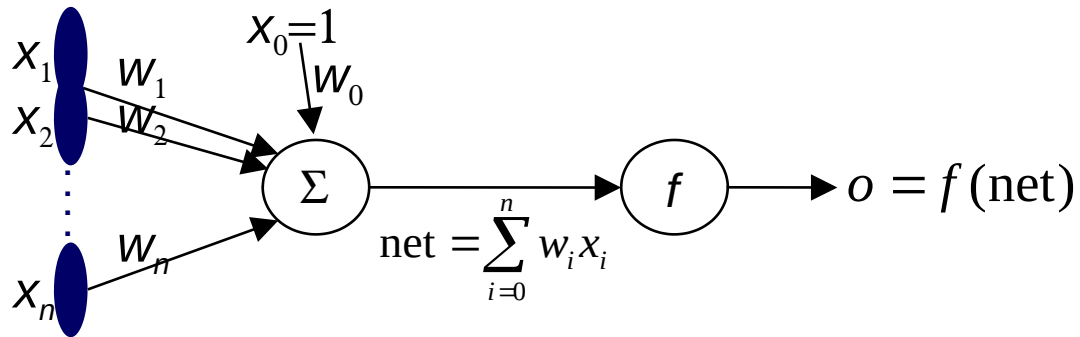
$$\frac{\partial E}{\partial w_i} = \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2$$

$$= \frac{1}{2} \sum_{d \in D} \frac{\partial}{\partial w_i} (t_d - o_d)^2$$

$$= \frac{1}{2} \sum_{d \in D} 2(t_d - o_d) \frac{\partial}{\partial w_i} (t_d - o_d)$$

$$= \sum_{d \in D} (t_d - o_d) \left(- \frac{\partial o_d}{\partial w_i} \right)$$

$$= - \sum_{d \in D} (t_d - o_d) \frac{\partial o_d}{\partial \text{net}_d} \frac{\partial \text{net}_d}{\partial w_i}$$



net: linear combination
o (output): logistic function

$$\frac{\partial o_d}{\partial \text{net}_d} = \frac{\partial f(\text{net}_d)}{\partial \text{net}_d} = f(\text{net}_d)(1 - f(\text{net}_d)) = o_d(1 - o_d)$$

$$\frac{\partial \text{net}_d}{\partial w_i} = \frac{\partial (w \cdot x_d)}{\partial w_i} = x_{i,d}$$

$$\frac{\partial E}{\partial w_i} = - \sum_{d \in D} (t_d - o_d) o_d (1 - o_d) x_{i,d}$$

... Incremental Version

- Batch gradient descent for a single Sigmoid unit

$$E_D = \frac{1}{2} \sum_{d \in D} (t_d - o_d)^2 \quad \frac{\partial E_D}{\partial w_i} = - \sum_{d \in D} (t_d - o_d) o_d (1 - o_d) x_{i,d}$$

- Stochastic approximation

$$E_d = \frac{1}{2} (t_d - o_d)^2 \quad \frac{\partial E_d}{\partial w_i} = - (t_d - o_d) o_d (1 - o_d) x_{i,d}$$

Backpropagation procedure


- Create FFnet
 - ◆ n_i inputs, n_o output units, n hidden layers
 - ◆ Define error by considering *all* output units
- Train the net by propagating errors backwards from output units
 - ◆ First output units
 - ◆ Then hidden units
- Notation: $x_{j,i}$ is input from unit i to unit j
 $w_{j,i}$ is the corresponding weight
- Note: various termination conditions
 - ◆ error
 - ◆ # iterations,...
- Issues of under/over fitting, etc.

Backpropagation (stochastic case)

- Initialize all weights to small random numbers
- Repeat
 - For each training example
 1. Input the training example to the network and compute the network outputs
 2. For each output unit k
$$\delta_k \leftarrow o_k (1 - o_k) (t_k - o_k)$$
 3. For each hidden unit h
$$\delta_h \leftarrow o_h (1 - o_h) \sum_{k \in \text{outputs}} w_{k,h} \delta_k$$
 4. Update each network weight $w_{j,i}$
$$w_{j,i} \leftarrow w_{j,i} + \Delta w_{j,i}$$
where $\Delta w_{j,i} = \eta \delta_j x_{j,i}$

Errors propagate backwards

$$\delta_3 = o_3(1 - o_3)(t_3 - o_3)$$



$w_{1,7}$ updated based on δ_1 and $x_{1,7}$

- Same process repeats if we have more layers

Properties of Backpropagation

- Easily generalized to arbitrary directed acyclic graphs
 - Backpropagate errors through the different layers
- Training is slow but applying it to networks after training is fast

Convergence of Backpropagation


● Convergence

- Training can take thousands of iterations → slow!
 - ♦ Gradient descent over entire network weight vector
 - ♦ Speed up using small initial values of weights:
 - Linear response initially
- Generally will find local minimum
 - ♦ Typically can find good approximation to global minimum
- Solutions to local minimum trap problem
 - ♦ Stochastic gradient descent
 - ♦ Can run multiple times – Over different initial weights
 - ♦ Committee of networks
 - ♦ Can modify to find better approximation to global minimum
 - include weight momentum α
$$\Delta w_{ij}(t_n) = \eta \delta_j x_{ij} + \alpha \Delta w_{ij}(t_{n-1})$$
 - Momentum avoids local max/min and plateaus

Example of learning a simple function


- Learn to recognize 8 simple inputs
 - ◆ Interest in how to interpret hidden units
 - ◆ System learns binary representation!
- Trained with
 - ◆ initial w_i between -0.1 , $+0.1$,
 - ◆ $\eta=0.3$
- 5000 iterations (most change in first 50%)
- Target output values:
 - ◆ 0.1 for 0
 - ◆ 0.9 for 1

Hidden layer representations



Input		Hidden values		Output
10000000	→		→	10000000
01000000	→		→	01000000
00100000	→		→	00100000
00010000	→	? ? ?	→	00010000
00001000	→		→	00001000
00000100	→		→	00000100
00000010	→		→	00000010
00000001	→		→	00000001

Hidden layer representations



Input		Hidden values		Output
10000000	→	.89 .04 .08	→	10000000
01000000	→	.01 .11 .88	→	01000000
00100000	→	.01 .97 .27	→	00100000
00010000	→	.99 .97 .71	→	00010000
00001000	→	.03 .05 .02	→	00001000
00000100	→	.22 .99 .99	→	00000100
00000010	→	.80 .01 .98	→	00000010
00000001	→	.60 .94 .01	→	00000001

Example of head/face recognition

- Task: recognize faces from sample of
 - 20 people in 32 poses
 - Choose output of 4 values for direction of gaze
 - 120x128 images (256 gray levels)
- Can compute many functions
 - Identity/direction of face (used in book)/...
- Design issues
 - Input encoding (pixels/features/?)
 - Reduced image encoding (30x32)
 - Output encoding (1 or 4 values?)
 - Convergence to 0.1/0.9 and not 0/1
 - Network structure (1 layer of 3 hidden units)
 - Algorithm parameters
 - $\eta = 0.3$; $\alpha = 0.3$; stochastic descent method
- Training/validation sets
- Results: 90% accurate for head pose


Some issues with ANNs

- Interpretation of hidden units
 - ◆ Hidden units “discover” new patterns/regularities
 - ◆ Often difficult to interpret
- Overfitting
- Expressiveness
 - ⋈ Generalization to different classes of functions


Dealing with overfitting

- Complex decision surface
- Divide sample into
 - ◆ Training set
 - ◆ Validation set
- Solutions
 - ◆ Return to weight set occurring near minimum over validation set
 - ◆ Prevent weights from becoming too large
 - ◆ Reduce weights by (small) proportionate amount at each iteration

Training vs. Validation



Effect of hidden units



Expressiveness

- Every Boolean function can be represented by network with a single hidden layer
 - Create 1 hidden unit for each possible input
 - Create OR-gate at output unit
 - *but* might require exponential (in number of inputs) hidden units
- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer (Cybenko et al '89)
 - Hidden layer of sigmoid functions
 - Output layer of linear functions
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers (Cybenko '88)
 - Sigmoid units in both hidden layers
 - Output layer of linear functions

Extension of ANNs

- Many possible variations
 - ◆ Alternative error functions
 - Penalize large weights
 - Add weighted sum of squares of weights to error term
 - ◆ Structure of network
 - Start with small network, and grow
 - Start with large network and diminish
- Use other learning algorithms to learn weights

Extensions of ANNs

- Recurrent networks
 - Example of time series
 - ♦ Would like to have representation of behavior at $t+1$ from arbitrary past intervals (no set number)
 - ♦ Idea of simple recurrent network
 - hidden units that have feedback to inputs
- Dynamically growing and shrinking networks

Inductive bias of Backpropagation

- Smooth interpolation between data points


Summary

- Practical method for learning continuous functions over continuous and discrete attributes
- Robust to noise
- Slow to train but fast afterwards
- Gradient descent search over space of weights
- Overfitting can be a problem
- Hidden layers can invent new features

Logistic function (Logit function)

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

This term lies in $[0, \infty)$



- $\sigma(z)$ is always bounded between $[0, 1]$ (a nice property),
- as z increases $\sigma(z)$ approaches 1,
- as z decreases $\sigma(z)$ approaches 0.

Segway: Logistic regression

- Logistic regression is often used because the relationship between the dependent discrete variable and a predictor is non-linear
- **Example:** the probability of heart disease changes very little with a ten-point difference among people with low-blood pressure, but a ten point change can mean a drastic change in the probability of heart disease in people with high blood-pressure.

Logistic regression

Learn a function to map X values to Y given data

$$(X^1, Y^1), \dots, (X^N, Y^N)$$

$$f: X \rightarrow Y$$

X can be continuous or discrete

Discrete

The function we try to learn is $P(Y|X)$

Logistic regression (Classification)

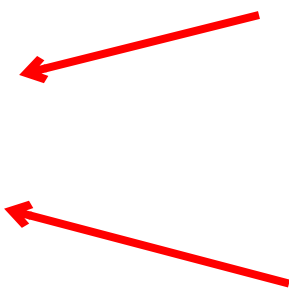
Classification

$$1 < \frac{P(Y = 0|X)}{P(Y = 1|X)}$$



If this holds $Y=0$ is more probable than $Y=1$ given X


Classification

$$1 < \frac{P(Y = 0|X)}{P(Y = 1|X)}$$



Take log both sides

Classification rule: if this holds $Y=0$

Logistic Function (Logit function)



Logistic regression is a linear classifier



Learn parameters using
sigmoid unit training
(gradient descent)

Thank You!

